

# Next-Day Stock Status Prediction in Fresh Retail

Anurag Thakur

2023IMG-008    Samyak Choudhary

2023IMG-044    Vibhor Kumar

2023IMG-054

Department of Information Technology  
Bachelor of Technology Project (BTP)

August 2026

# Problem Statement & Objective Metrics

## Operational Challenge in Fresh Retail

- Perishable fresh items suffer from **rapid intra-day demand shocks** and supply replenishment lags.
- Unpredicted stockouts cause direct financial revenue loss, stock spoilage, and customer churn.

## Predictive Target Formulation

- Forecast a **24-dim binary hourly status vector**:  
 $\mathbf{Y} = [y_1, y_2, \dots, y_{24}]^T \in \{0, 1\}^{24}$
- Covers active operational store hours (**6 AM to 10 PM**).
- $y_h = 1$  indicates SKU stockout during hour  $h$ .

## Key Objective Performance Metrics

- **Hour-Level F1-Score** (Primary Metric):  
Evaluates precision vs. recall on positive stockout hours:  $F1 = \frac{2 \cdot P \cdot R}{P + R}$
- **Stockout Recall**: Measures percentage of true operational stockout hours correctly detected ( $R = \frac{TP}{TP + FN}$ ).

## Taxonomy of Time-Series Forecasting Literature

- **Recurrent Sequence Memory** (Hochreiter & Schmidhuber, 1997): *LSTM* maintains persistent cell state memory  $C_t$  for temporal sequence continuity.
- **Deep Time-Series Transformers** (Lim et al., 2021 [TFT]; Zeng et al., 2023 [PatchTST]): Self-attention mechanisms suffer from **overparameterization** ( $> 280k$  params) on short sequences.
- **Linear Series Decomposition** (Zeng et al., 2023 [DLinear]; Challu et al., 2023 [N-HiTS]): Moving-average trend decomposition provides parameter-efficient baselines ( $< 15k$  params).
- **Explainable AI in Retail** (Sundararajan et al., 2017 [IG]; Lundberg & Lee, 2017 [SHAP]): Attribution methods reveal feature-level importance to guide neural architecture design.
- **FreshRetail Benchmark Studies** (Same Dataset): Recent domain-specific works like "*A Coupled Hierarchical Architecture for Multi-Granularity Demand Forecasting*" and "*When Less Data Is Better: Eliminating Stockout-Induced Bias in Choice Estimation*" use this dataset to model complex stockout biases.

# FreshRetailNet-50K Dataset Overview

## Dataset Scale & Scope

- **Volume:** 50,000 store-product pairs across 898 stores.
- **Timeframe:** 90-day series of highly granular hourly data.
- **Products:** 865 diverse perishable SKUs.

## Rich Feature Set

- **Sales & Stock:** Hourly sale amounts & exact stock status.
- **Events:** Discount rates, holiday flags, and activities.
- **Weather:** Precipitation, temp, humidity, and wind.

## Why We Selected This Dataset

- **Pioneering Benchmark:** It is the first large-scale public dataset specifically designed for **censored demand estimation** (where out-of-stock items hide true customer demand) in the fresh retail domain.
- **Perishable Goods Focus:** Perfectly captures the unique, rapid intraday volatility of short-shelf-life items, containing ~20% organically occurring stockout data.

# Data Selection & Splitting Strategy

## Unit of Analysis & Filtering

- **Target Grain:** Predicted at the level of `city_id + store_id + product_id`.
- **Filtering Logic:** Fresh retail datasets contain thousands of dormant items. We select the **Top 15 SKUs** to focus on high-impact, stockout-prone goods.

## SKU Statistical Ranking Formula

$$\text{Score} = 0.40 \times \text{Sales} + 0.25 \times \text{Stockout Hours} + 0.20 \times \text{Stockout Days} + 0.15 \times \text{Sales Std}$$

## Train/Validation Split

- **Chronological Split:** Strict temporal separation to prevent data leakage.
- **Training Set:** Historical data up to  $T_{\max} - 15$  days ( $\sim 2.5$  months).
- **Validation Set:** Final 15 days of data.
- **Standardization:** `StandardScaler` fitted strictly on the training set to prevent lookahead bias.

# Hourly vs. Daily Forecasting Paradigm


## Traditional Daily Aggregation Paradigm

- **Input:**  $L = 14$  days of historical daily metrics.
- **Output:** 1 scalar value (Total expected daily sales or binary stockout flag).
- **Loss of Information:** Store managers cannot determine *when* replenishment must occur.

## Our 24-Hour Hourly Binary Paradigm

- **Input:**  $L = 14$  days of SKU-store daily feature vectors.
- **Output:** 24 binary probabilities  $\hat{\mathbf{Y}} \in [0, 1]^{24}$ .
- **Actionable Insight:** Identifies peak risk hours (e.g. 11 AM lunch rush & 6 PM evening peak) for proactive restocking.


# Model Family Comparison: Simple LSTMs Beat Heavy Transformers




## Key Empirical Findings

- **Parameter Efficiency:** Normal LSTMs (5.9k–44.9k params) reach **0.7852–0.7897 F1-Score**, matching or beating heavy Transformers.
- **Transformer Overparameterization:** PatchTST (285k params) & TFT (288k params) peak at **0.7540–0.7685 F1-Score**.
- **Why LSTMs Win:** Sequential state memory  $h_t$  naturally aligns with continuous hourly depletion better than permutation-invariant self-attention.

# ML Explainability: Integrated Gradients Analysis



Attribution heatmap: all 11 features & 14-day history (D-1 = most recent)




## Observation → Selection


Out of all 11 features, IG reveals the same 6 features consistently dominate across time steps and samples. Weather & activity features show near-zero attribution. This motivates our **feature reduction to 6 signals** and the dual-stream split between demand & inventory.



# Limitations of Baseline LSTM



Standard LSTM Cell Architecture




F1-Score across LSTM hidden size variants ( $h=8$  to  $h=96$ )

## Why Scaling Alone Doesn't Help


- **F1 plateau** at  $\sim 0.789$  regardless of hidden size — more params yield no gain.
- Demand & inventory share hidden state  $h_t \Rightarrow$  **feature cross-talk**.
- IG analysis (prev. slide) proves these signals need **dedicated streams**.

# Our Fix: Dual-Stream Inventory Shortcut LSTM



Dual-Stream: Demand LSTM + Inventory LSTM + Day-14 Shortcut

LSTM Architecture Evolution: Baseline → Dual-Stream Variants



F1, Recall & MAE: all dual-stream variants vs. baseline LSTM

## Key Results

- **Inventory Shortcut:** F1 0.7926, MAE 3.97 hrs at just 4,600 params.
- Decoupled streams consistently beat all baseline LSTM sizes.

# Exploring Linear Decomposition: DLinear Models

## Why Explore Linear Decomposition?

DLinear (Zeng et al., 2023) demonstrated that single-layer linear models decomposing time series into Trend and Seasonal components can outperform complex deep networks.

## DLinear Series Decomposition


$$X_{\text{trend}} = \text{AvgPool}(X), \quad X_{\text{seasonal}} = X - X_{\text{trend}}$$

$$Y_{\text{pred}} = W_{\text{trend}} X_{\text{trend}} + W_{\text{seasonal}} X_{\text{seasonal}}$$

## Standard DLinear Performance

- **Parameters:** 6,768 parameters.
- **Validation F1-Score: 0.7932**
  - Outperforms Normal LSTMs (**0.7897**)
  - Outperforms Enhanced LSTMs (**0.7926**)
  - Outperforms Transformers (**0.7540**)
- **Stockout Recall: 87.44%.**

# Architectural Innovation 2: Linear Changes in DLinear




## Hourly Slot-Specific Linear DLinear

- Uses **24 dedicated linear heads** ( $W_1, \dots, W_{24}$ ) for each operational hour.
- Resolves inter-hour gradient conflict between morning and evening rush hours.
- Boosts F1 to **0.8058** (+0.0161 over LSTM).

## Multi-Kernel Linear DLinear

- Multi-scale moving average trend pooling ( $k \in \{3, 5, 7\}$ ).
- Reaches **0.8131 F1-Score** with pure linear decomposition (13.5k params).

# Architectural Innovation 3: Introducing Non-Linearity into DLinear



## Non-Linear Activation Bottleneck


Adds GELU activation bottlenecks between decomposition and output projections:

$$\hat{Y} = W_{\text{trend},2} \text{GELU}(W_{\text{trend},1} X_{\text{trend}}) + \text{Seasonal}$$

## Breakthrough Results

- **Non-Linear GELU DLinear:** Reaches **91.37% Stockout Recall** (Highest in Repo!) and **0.8083 F1**.
- **Multi-Kernel Non-Linear DLinear:** Reaches **0.8156 F1-Score** (New Peak F1 & lowest 3.75h MAE).

# Architectural Schematic: Multi-Kernel & Slot DLinear



## Core Mechanism

- **Trend Stream:** Multi-kernel moving average pooling ( $k \in \{3, 5, 7\}$ ) + GELU bottleneck.
- **Seasonal Stream:** High-frequency residual stockout shocks ( $X - X_{trend}$ ).
- **24 Slot Heads:**  $W_1, \dots, W_{24}$  dedicated projections resolve inter-hour gradient conflicts.

# Master Empirical Validation Leaderboard

| Rank        | Architecture Name                      | Category             | Params  | F1-Score      | Recall        | MAE (h)     |
|-------------|----------------------------------------|----------------------|---------|---------------|---------------|-------------|
| 1           | <b>Multi-Kernel Non-Linear DLinear</b> | Non-Linear DLinear   | 42,339  | <b>0.8156</b> | 87.13%        | <b>3.75</b> |
| 2           | <b>Multi-Kernel Linear DLinear</b>     | Linear (Multi-Scale) | 13,539  | <b>0.8131</b> | 85.22%        | 3.86        |
| 3           | <b>Non-Linear GELU DLinear</b>         | Non-Linear DLinear   | 21,168  | <b>0.8083</b> | <b>91.37%</b> | 4.20        |
| 4           | Hourly Slot-Specific DLinear           | Linear (24 Heads)    | 6,744   | <b>0.8058</b> | 88.02%        | 4.23        |
| 5           | Standard Linear DLinear                | Linear (Decomp)      | 6,768   | <b>0.7932</b> | 87.44%        | 4.28        |
| 6           | Dual-Stream Inventory Shortcut         | LSTM Family          | 4,600   | <b>0.7926</b> | 80.79%        | 3.96        |
| <b>Base</b> | <b>Best Normal LSTM (h=32)</b>         | Normal LSTM          | 5,912   | <b>0.7897</b> | 81.29%        | 4.31        |
| 7           | Feature-Reduced Standard LSTM          | Normal LSTM          | 19,992  | 0.7880        | 81.60%        | 4.58        |
| 8           | Baseline Standard LSTM (h=96)          | Normal LSTM          | 44,952  | 0.7852        | 80.40%        | 4.67        |
| 9           | Vanilla Simple RNN (h=64)              | RNN Family           | 14,744  | 0.7750        | 81.02%        | 4.23        |
| 10          | PatchTST Patch Linear                  | Transformer          | 285,624 | 0.7540        | 81.90%        | 4.16        |

# Conclusion & Operational Deployment Takeaways

## Key Architectural Takeaways

- 1 **Non-Linearity Boosts DLinear:** GELU bottlenecks boost DLinear F1 to **0.8083** and Recall to a record **91.37%**.
- 2 **New All-Time Peak F1:** Multi-Kernel Non-Linear DLinear sets the highest F1-Score of **0.8156** and lowest MAE of **3.75 hours**.
- 3 **Feature Decoupling & Slot Heads:** Resolves inter-hour gradient noise and temporal cross-talk.

## Deployment Recommendations

- **Primary Production Deployment:** **Multi-Kernel Non-Linear DLinear** (0.8156 F1, 3.75h MAE, 42.3k params).
- **Maximum Risk Mitigation:** **Non-Linear GELU DLinear** (91.37% Stockout Recall).
- **Future Scope:** Incorporate real-time intraday point-of-sale stream updates for dynamic threshold adjustment.